CTO REPORT | PROMPT SECURITY

7d of activity under policy

What Prompt Security observed, what we intervened on, and where SentinelOne takes the insight next.

Window 2026-04-15 to 2026-04-22 UTC

`Data source  dataSource.name = 'Prompt Security'`

SentinelOne Pre-Sales Engineering

Generated 2026-04-22

Executive summary

37,094	15,922	47.5%	40
Total events (7d)	Blocked events	Policy intervention rate	Distinct users

Three things for the CTO to act on

Block 42.9%, bypass 0.135%.

Intervention rate is 47.5%. The policy is doing heavy work. Confirm the traffic mix justifies this posture before moving anything to enforce.

One user drove 96.3% of volume.

drjones@flybynight.ru drove 96.3% of all events. Resolve this account before treating any other number as a baseline.

Bypass is the signal to watch.

Bypass is near-zero (0.13%). Users are accepting policy decisions. The cleanest possible trust signal.

Policy action mix

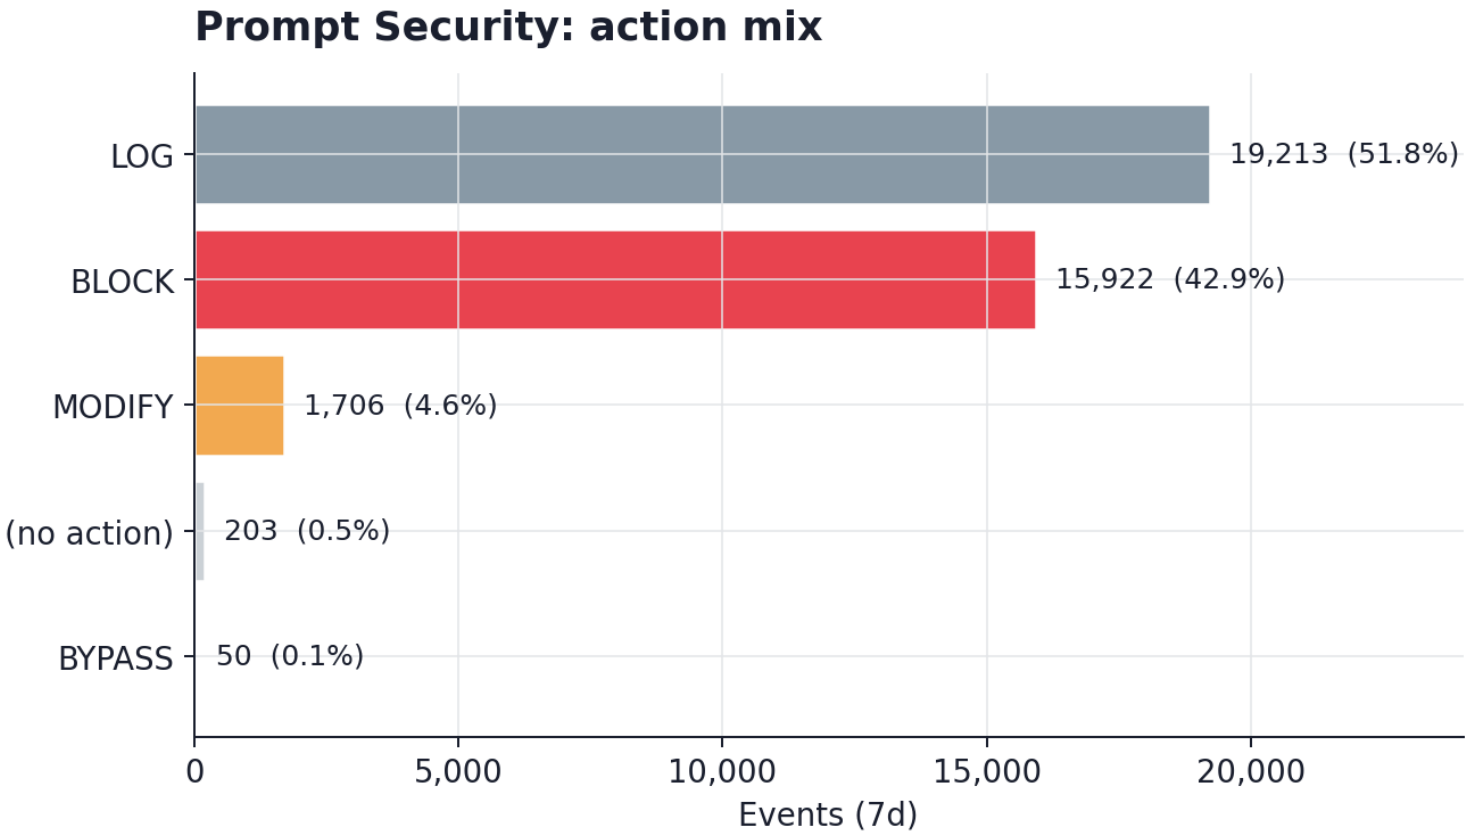

What to notice

- BLOCK at 42.9%. Events the policy refused outright; exfil, compliance, jailbreak attempts live here.
- MODIFY at 4.60%. Payload rewritten in place (eg. PII redacted) without stopping the user.
- BYPASS at 0.135%. Users almost never override. Policy is trusted.

Volume and timing

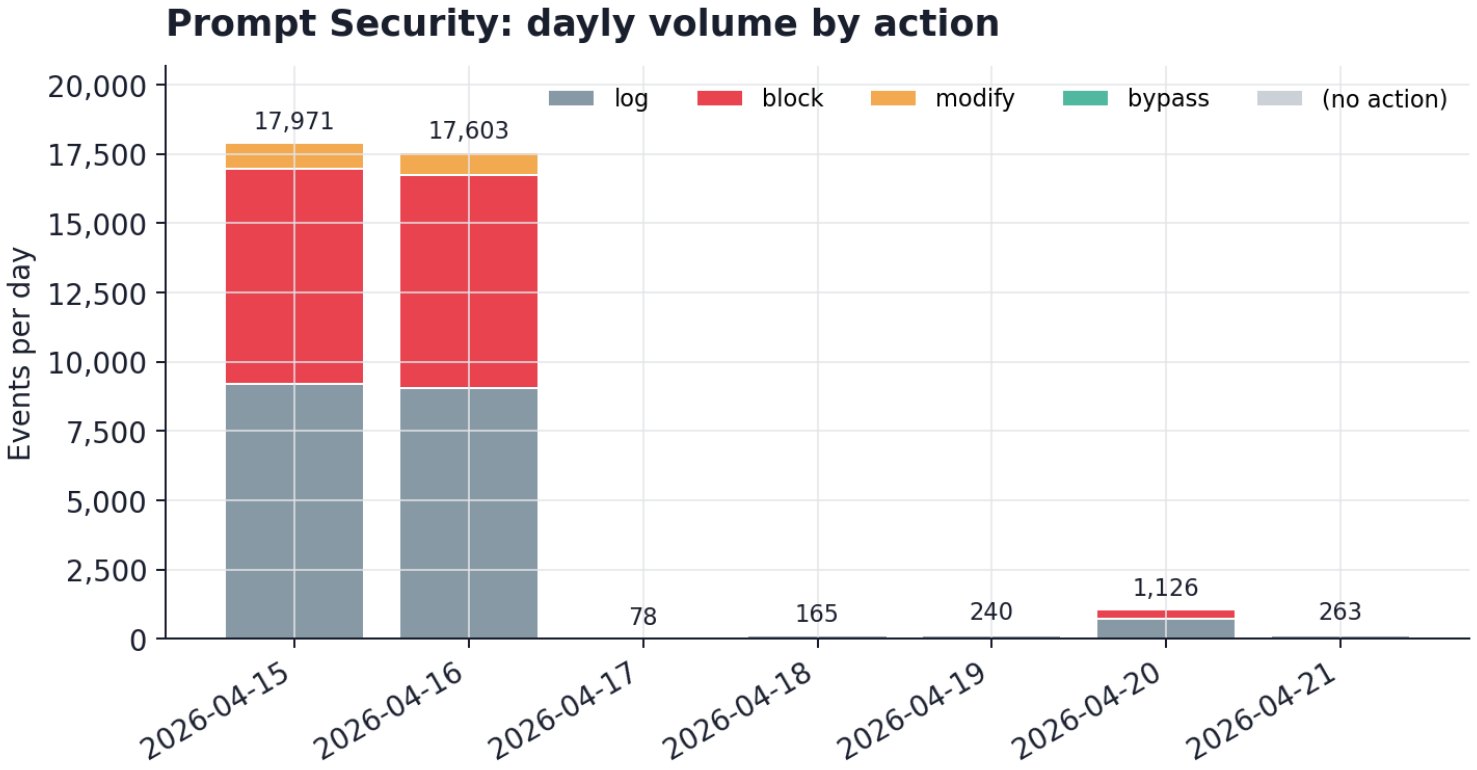

What to notice

- Uneven but not pathological: busiest slice (2026-04-15) carried 48.0% of the window.
- Dominant action across the window: 'log'. Watch the ratio to the next action as an early policy-drift signal.
- Policy decisions are stable under load across the window. That is the precondition for promoting any log-only policy to enforce.

Who is driving the traffic

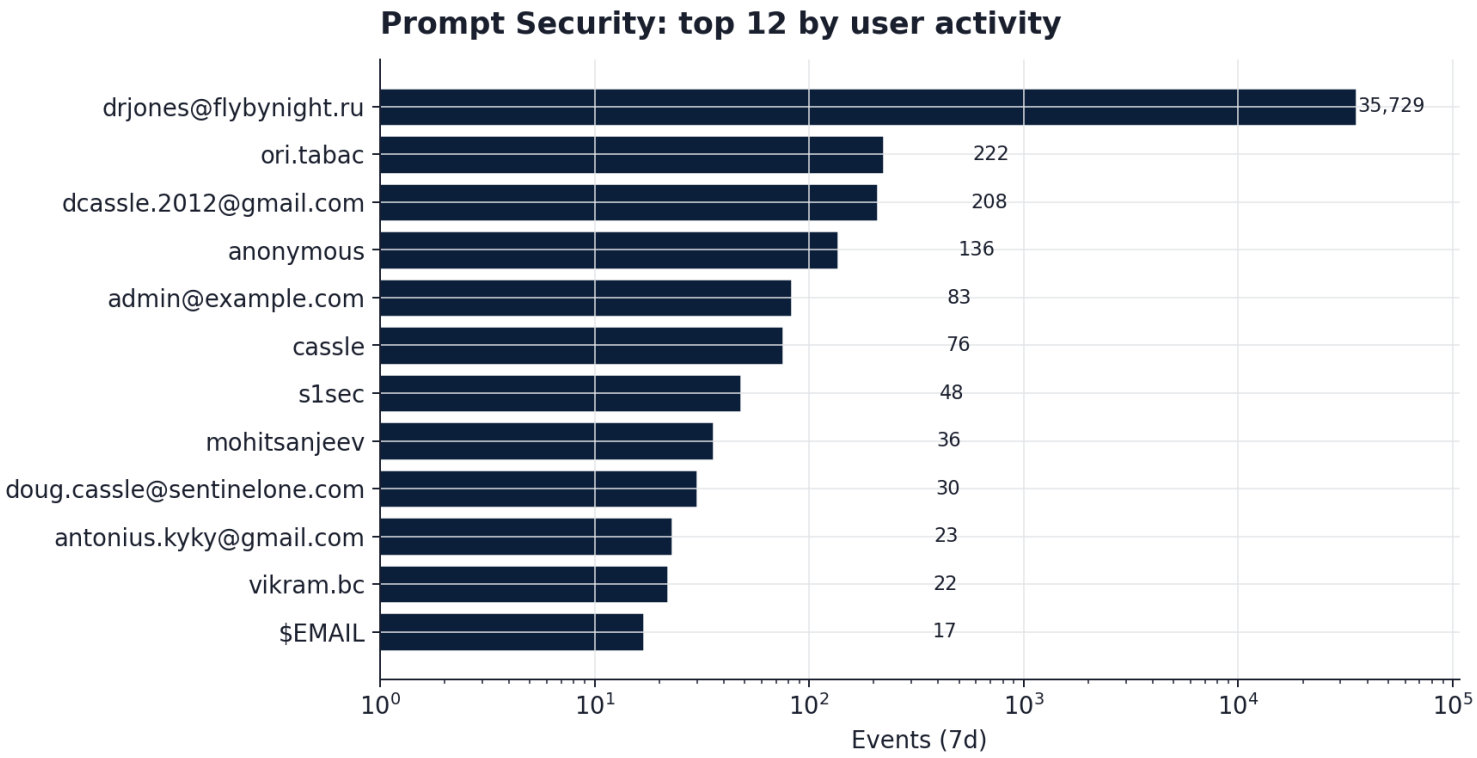

What to notice

- drjones@flybynight.ru = 35,729 events (96.3% of volume).
- The rest of the cohort lives in the 1 to 222 range. That is where normal users behaviour hides.
- Log scale used deliberately so a dominant principal does not flatten the long tail.

Highest-risk prompt senders

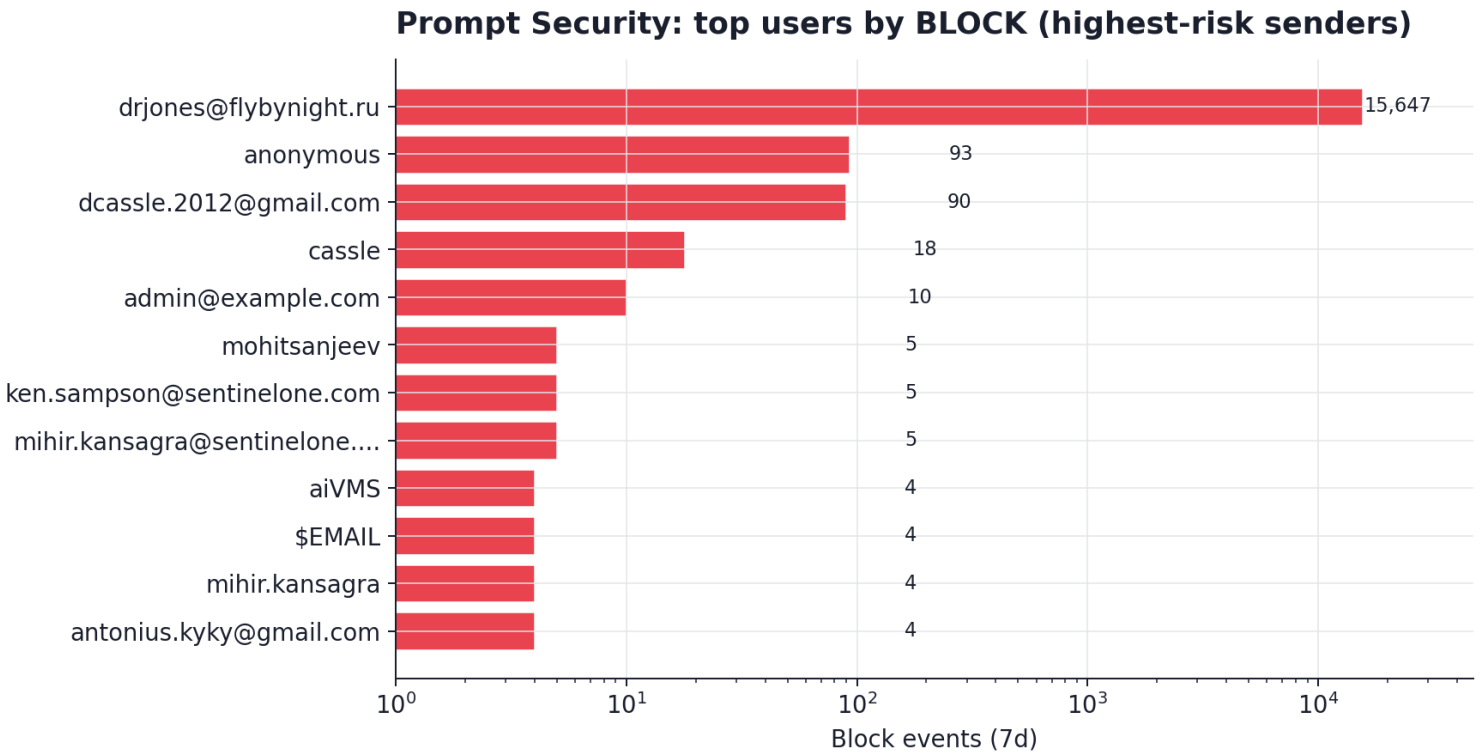

What to notice

- drjones@flybynight.ru blocks = 15,647. That is 98.3% of all blocks in the window.
- The rest of the blocked cohort sits between 1 and 93 blocks per user. Policy tuning or 1:1 coaching has the highest ROI here.
- Anyone labelled 'anonymous' or 'null' in the chart warrants a tagging pass so blocks can be attributed.

Who is overriding guardrails

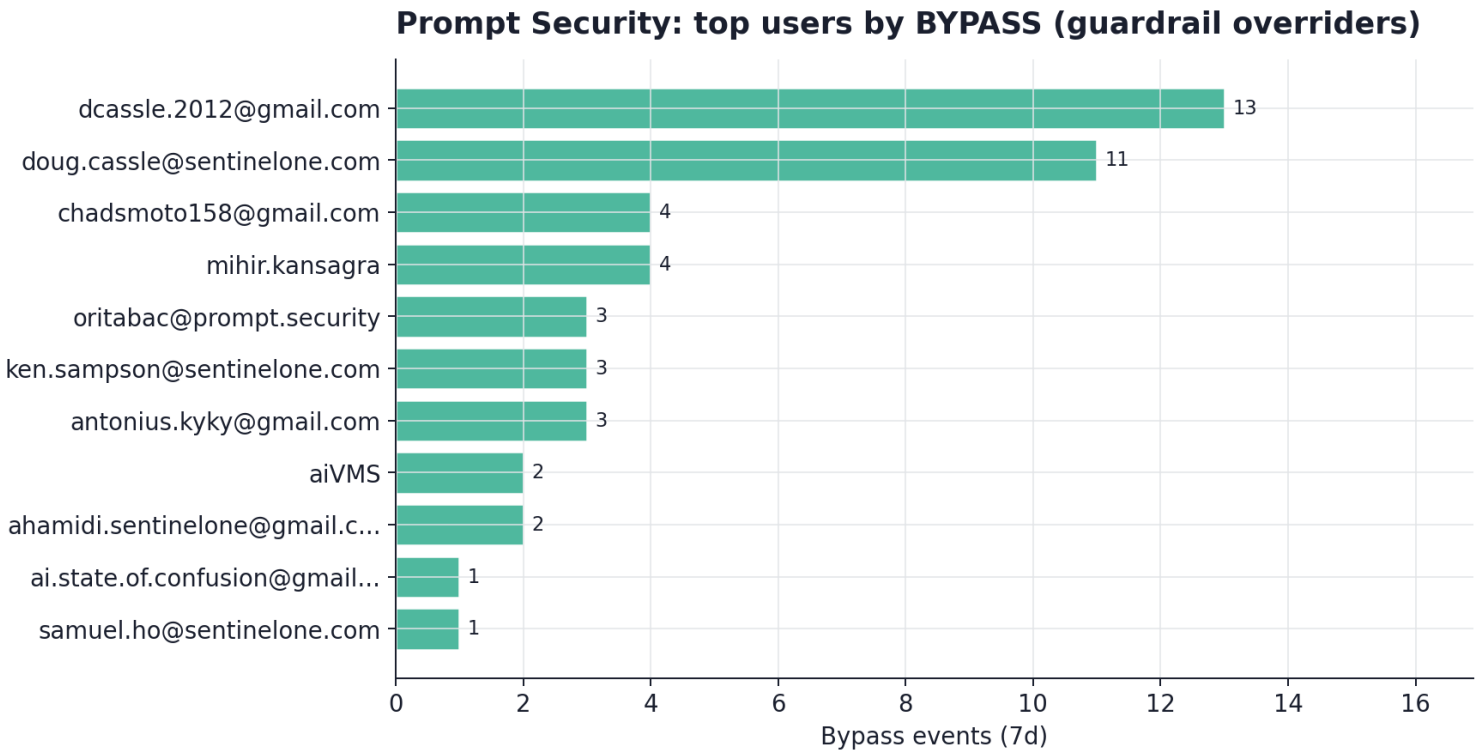

What to notice

- 50 bypasses across the window. Low-volume, high-signal.
- Top overriders: dcassle.2012@gmail.com, doug.cassle@sentinelone.com. Cross-check endpoint posture for each.
- Bypass is near-zero (0.13%). Users are accepting policy decisions. The cleanest possible trust signal.

Per-principal action mix

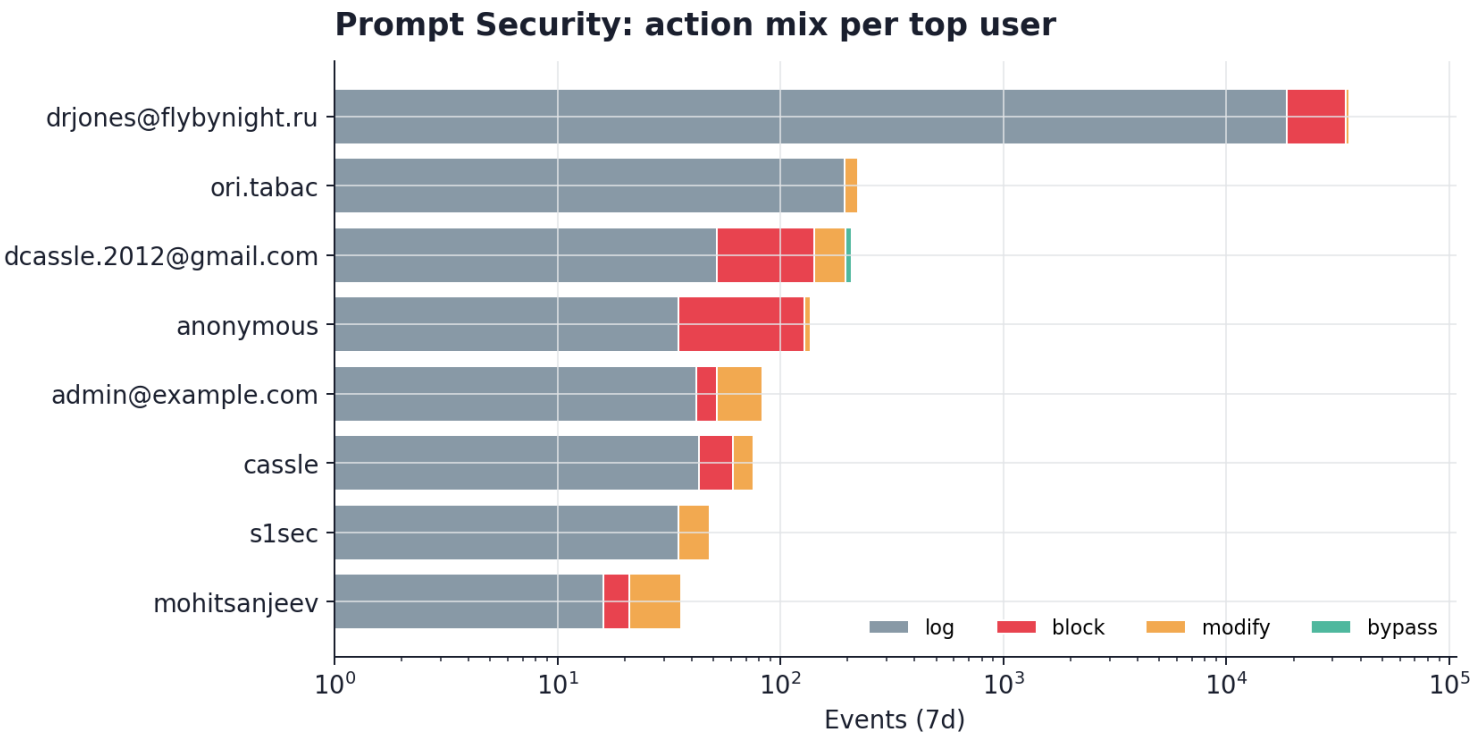

What to notice

- Stacked action mix per top users. This is the fastest way to spot a mixed-use profile (blocked and allowed inside one session).
- oritabac@prompt.security shows the clearest mixed-use profile (block:3, log:3, bypass:3, modify:2). Likely a real analyst probing the guardrails.
- Use this chart to pick principals for policy feedback loops before the next re-run.

Where Prompt Security sits in the broader tenant

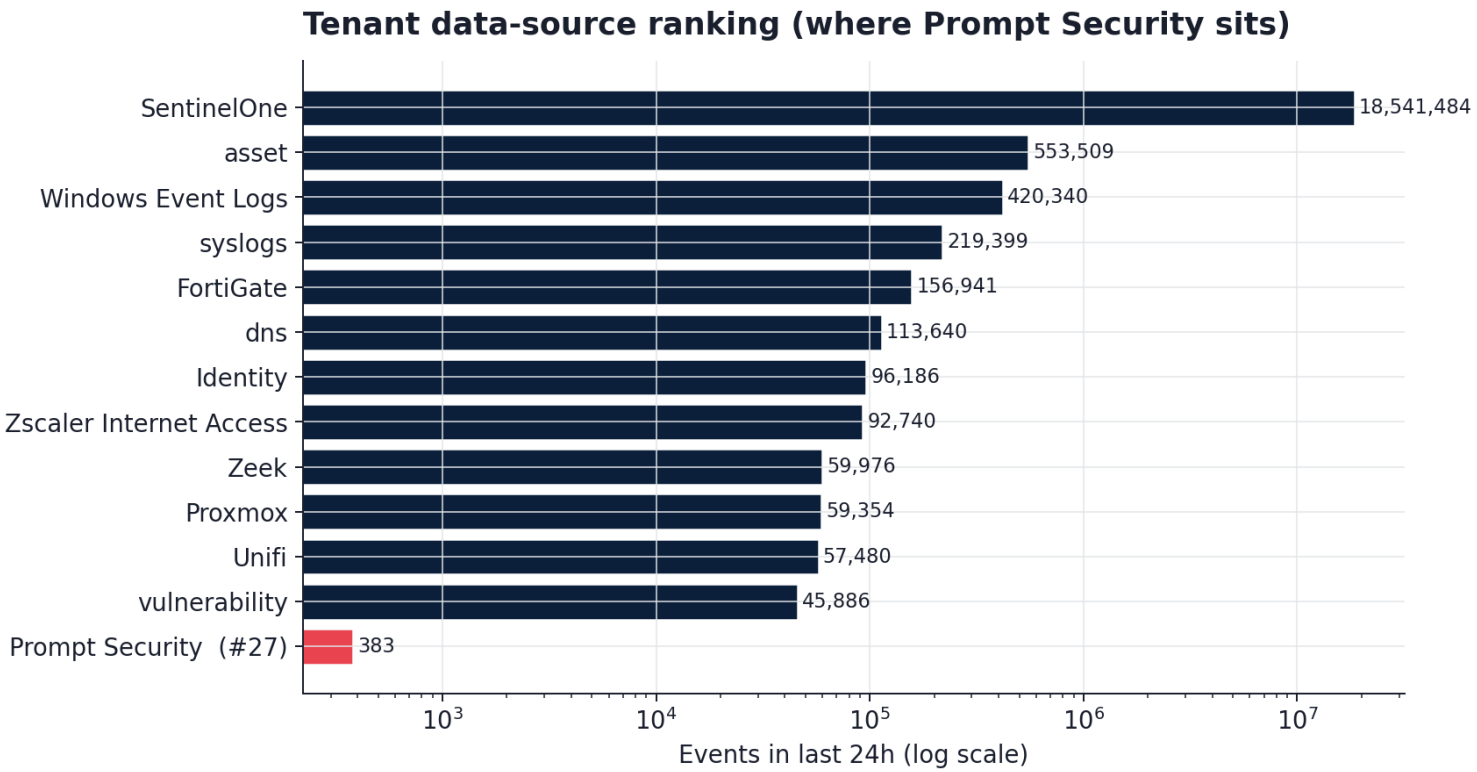

What to notice

- In the last 24 hours, Prompt Security ranks #27 among the data sources landing in this tenant.
- Same data lake as SentinelOne EDR, identity, network, and cloud telemetry. Correlation is a PowerQuery join away.
- As usage grows, Prompt Security's share of daily events climbs in rank with zero ingestion changes.

Recommendations, prioritised

24 HOURS

Resolve drjones@flybynight.ru

A single user drove 96.3% of the window. Pivot into EDR / identity from that session to establish harness vs. shared-account vs. abuse. Close the single largest risk in this report.

THIS WEEK

Promote log-only policies to enforce

Block 42.9% and bypass 0.135%. That combination is as clean a trust signal as we get. Walk each log-only policy and flip the ones with zero false positives.

THIS MONTH

Correlate bypass with endpoint posture

STAR rule: Prompt Security bypass THEN a SentinelOne alert of severity medium+ on the same principal within 60 minutes. One rule, two data sources, one unified alert.

THIS MONTH

Enrich ingestion where it is shallow

Today, Prompt Security lands with a subset of fields populated. Add prompt / payload / content where possible; that unlocks content-based detections and Purple AI hunts across the same data lake.

NEXT STEPS

Resolve drjones@flybynight.ru, promote log-only policy, enrich ingestion.

Everything in this deck is reproducible from a single JSON artefact.
Re-run any time the workload shifts.

Thank you. Questions?